

AI Hiring Vendor Evaluation Checklist

The questions to ask any AI interview or assessment vendor — and the answers that should end the meeting.

Transparency

- [] For any score, can we see the evidence trail — transcript, rubric anchor, reasoning?
- [] Can the candidate-facing experience be demoed end to end?
- [] Is there a documented model/system card describing what the AI does and doesn't do?

Fairness engineering

- [] What is redacted before scoring (names, demographic signals, accent markers)?
- [] Can they show adverse-impact monitoring (e.g., four-fifths analysis) — actual reports, not policy statements?
- [] Is scoring validated against calibrated human expert scores? How often is drift checked?

Human oversight

- [] Which decisions does a human make, and where is that recorded?
- [] What happens when the AI is uncertain — flagged for review, or silently rounded?
- [] Can a candidate request human review where regulation requires it?

Data & compliance

- [] Data processing agreement, retention policy, and deletion path in writing
- [] Candidate notice and consent flows appropriate to your jurisdictions (NYC LL144, EU AI Act, etc.)
- [] Where is data stored and who can access it?

Red flags — end the meeting

- Facial-expression, tone-of-voice, or personality inference 'analysis'
- Scores with no visible reasoning or transcript
- Claims of 'eliminating' bias (serious vendors reduce and measure)
- No answer to 'show me one candidate's full evidence trail'

© 2026 L3XY AI — free to use and share within your organization. L3XY runs structured AI interviews that turn every candidate into verified hiring signals. Learn more at www.l3xy.ai