

Technical Memo: AI Interview Consent — Jurisdiction Capability (Current State)

Prepared for: Outside counsel **Date:** July 7, 2026 **Scope:** Factual description of the Lexy platform's current consent/disclosure implementation. No legal conclusions or proposed consent language. All quoted text is verbatim from the running system.

1. Current state — what the candidate sees before consenting

There is **one fixed AI-interview disclosure, shown identically to every candidate everywhere**, defined in a single source file (`artifacts/api-server/src/lib/ai-consent.ts`) under version string `ai-interview-2026-06` . It is served to two consent surfaces (the logged-in candidate portal consent page, and the in-room consent gate on public interview links); both render the same structured object. Verbatim content:

Intended use:

"Lexy will conduct a structured video / voice interview with you and produce a summary of your answers for the recruiting team. The recruiter, not Lexy, decides whether to advance you in the process."

Model providers disclosed: "OpenAI (GPT-4o family)", "Anthropic (Claude 3.5 Sonnet)" **Decision maker disclosed as:** "human-in-the-loop"

Evaluated traits (shown as a list):

- *"Role-relevant skills and experience as expressed in your resume"*
- *"Role-relevant skills and experience as expressed during the interview"*
- *"Clarity and structure of your spoken answers"*
- *"Demonstrated reasoning on the questions asked"*

Explicitly NOT evaluated (shown as a list):

- *"Race, ethnicity, or national origin (beyond work authorisation, which you self-report)"*

- *"Age, religion, sexual orientation"*
- *"Disability, family or marital status"*
- *"Gender identity"*
- *"Accent, pitch, or speech characteristics as a proxy for ability"*
- *"Facial expressions or any video-based personality inference"*

Candidate rights:

- *"You may withdraw consent at any time from your candidate portal."*
- *"You may request deletion of your interview recording and transcript within 30 days under the Illinois AIVI Act, or at any time under GDPR Article 17 / CCPA."*
- *"You may request a copy of your data."*

Retention:

"Interview recordings are retained for the duration of the active hiring process plus 12 months for audit purposes, then deleted. Deletion requests are honored within statutory windows."

Biometric section — identifiers collected:

- *"A scan of your facial geometry and gaze / attention signals captured from your webcam during the proctored interview"*
- *"A recording of your voice and a video recording of the interview session"*

Biometric purpose:

"These biometric identifiers and biometric information are collected only to (a) confirm that it is you taking the interview, (b) maintain interview integrity (proctoring), and (c) create a record of your interview for the recruiting team to review. They are NEVER used to infer personality, emotion, demographics, or any protected characteristic."

Biometric retention schedule:

"Your biometric identifiers and biometric information are permanently destroyed when the initial purpose for collecting them has been satisfied — i.e. when the hiring process concludes — or within 1 year of your last interaction with Lexy, whichever occurs first, in accordance with the Illinois Biometric Information Privacy Act (740 ILCS 14). You may request earlier deletion at any time."

No sale/sharing:

"Lexy does not sell, lease, trade, or otherwise profit from your biometric data, and does not disclose it to any third party without your consent or as required by law."

Consent mechanics: the candidate must affirm **two separate checkboxes** — the AI-interview disclosure and a separate biometric release — both required as literal `true` values by the server. The consent record stores: candidate ID, version string, a **full snapshot of the disclosure text as shown** (so the exact copy is reconstructable even after later edits), timestamp, user-agent, IP, the separate biometric affirmation, and a revocation timestamp if later withdrawn (rows are never deleted). English only; no translations exist.

Storage/versioning: text lives in code (one constant + one function), versioned by the string `ai-interview-2026-06`. Changing the copy requires a code deployment and a manual version bump; a bump forces all candidates to re-consent.

2. Does any jurisdiction variance exist today?

For the AI-interview/biometric consent above: no. One disclosure, one version, every candidate, every location.

However, the platform already contains two OTHER consent/disclosure systems that DO vary by geography (facts counsel should know exist):

(a) Jurisdiction-aware pre-decision notices (separate from interview consent).

A governance layer classifies candidates/jobs into jurisdiction codes — `US-NY-NYC`, `US-CO`, `US-IL`, `EU` — from free-text location fields (`lib/governance/jurisdictions.ts`). It is deliberately over-inclusive: an empty/unknown location is treated as ALL four regulated jurisdictions; "New York" (state) is treated as NYC. A database table (`jurisdiction_disclosure_templates`) stores notice copy **keyed by (jurisdiction_code,**

language, template_key) with template keys `aedt_notice`, `co_pre_decision`, `il_aivi`, `eu_ai_act` — currently seeded with placeholder copy the code comments say "Legal will edit." Candidates' acknowledgements of these notices are recorded append-only (exact template + policy version IDs, IP, UA; the table has database triggers preventing UPDATE/DELETE). This system serves *notices*; it is **not wired into the interview consent gate**.

(b) Region-aware voluntary self-identification (demographics). The self-ID flow already swaps disclosure copy by region: EU/UK tenants get GDPR Article 9 special-category explicit-consent text; US/CA/AU tenants get OFCCP-style "Voluntary Self-Identification" text. Region is derived from the **tenant's (employer's) configured region**, not the candidate's location. Versioned separately (`self-id-2026-05`), snapshot on save.

3. Technical capability for variance — where a jurisdiction-keyed consent would plug in

- The consent text is rendered from a **single constant and a single function** (`CURRENT_AI_CONSENT_VERSION` / `getCurrentDisclosure()`), consumed by all consent surfaces and all enforcement gates. Replacing these with a jurisdiction-keyed lookup is a **localized change with moderate effort**: every surface would inherit the variant automatically.
 - The **classifier already exists** (§2a) and could supply the jurisdiction key; its "unknown location → all regulated jurisdictions" posture would default unknown candidates to the strictest applicable variant.
 - The **audit storage already supports variants**: the consent table snapshots the full disclosure as shown and records a free-form version string, so per-jurisdiction versions (e.g. `ai-interview-2026-06/US-IL`) would be self-documenting with no schema change.
 - The **database-template pattern already exists** (§2a's jurisdiction/language-keyed table with placeholder copy awaiting legal edit) and could be extended or mirrored for interview consent if counsel wants copy editable without code deployments.
 - What does NOT exist: any per-jurisdiction consent *logic* (e.g., differing checkbox requirements, cooling-off periods, guardian consent, or jurisdiction-specific revocation mechanics). Only the text-selection layer has precedent; behavioral differences would be new construction.
-

4. Is candidate location captured anywhere the system could key off?

Yes — several signals exist today, of varying reliability:

Signal	Where	Nature
Self-reported location	<code>candidates.location</code> (free text, e.g. "Chicago, IL")	Optional; candidate- or recruiter-entered; the jurisdiction classifier (§2a) is built to parse exactly this field
Job location	<code>jobs.location</code> (free text)	Used by the classifier alongside candidate location
IP address at consent	Stored in the consent record's capture context	Recorded for audit; not currently geo-resolved to a jurisdiction
Phone number	<code>candidates.phone</code> (free text)	Country code not parsed or used for geography
Work-authorization country	<code>sponsorship_country</code> on candidates	Self-reported during interviews; indicates authorization target, not residence
Employer region	<code>tenants.region</code>	Configured per employer; already drives the self-ID copy variant (§2b)

No geocoding/IP-geolocation service is integrated. The classifier's stated design position (in code comments): free-text heuristics are acceptable for compliance gating because it fails toward over-inclusion; a real geocoder would be added "when precision matters more."

5. Named regimes already referenced in the product's own code/notices (surface area only)

Factual inventory of regimes the system's disclosures, comments, or tables name today — listed for orientation, with no interpretation:

- **Illinois AIVI Act (820 ILCS 42)** — named in the disclosure's deletion-rights text and consent-table documentation
- **Illinois BIPA (740 ILCS 14)** — named verbatim in the biometric retention schedule shown to candidates

- **NYC Local Law 144 (§ 5-301)** — named in the notice system (§2a); `aedt_notice` template key seeded
 - **Colorado SB24-205** — named in the notice system; `co_pre_decision` template key and `US-C0` policy seeded
 - **EU AI Act (incl. Article 26(11))** — named in consent-table docs and notice system; `eu_ai_act` template key seeded
 - **GDPR** — Article 17 (deletion) named in the candidate-rights text; Article 9 (special categories) drives the EU/UK self-ID copy
 - **CCPA** — named in the candidate-rights deletion text
 - **OFCCP voluntary self-identification** — drives the US/CA/AU self-ID copy and the decoupled demographics store (k-anonymized aggregates only, never joined to recruiter views)
-

Summary for counsel

One fixed, English-only AI/biometric disclosure (version `ai-interview-2026-06`) is shown to all candidates globally, with full-text snapshot auditing and two-part affirmation. No jurisdiction branching exists in the interview consent flow today, but the platform already operates (i) a jurisdiction classifier over self-reported locations, (ii) a jurisdiction/language-keyed notice-template table awaiting legal copy, and (iii) a region-varied self-ID disclosure — so the technical pattern for per-jurisdiction consent text exists and the plug-in point is a single constant/function. Candidate location exists as free text plus consent-time IP (not geo-resolved); no geocoding is integrated. Behavioral (non-text) per-jurisdiction consent differences would be new construction.