

EU AI ACT PACK

AI System Card

Model inventory & intended use

Lexy AI Hiring Platform · Generated 2026-07-20

Factual description of the system as built. Not legal advice; assumptions flagged for counsel.

This document is Lexy’s lightweight AI System Card. It maps to the information required by the EU AI Act for a high-risk AI system (Annex IV, “Technical documentation”) and serves as the public-facing explanation of how Lexy uses AI to evaluate candidates.

It deliberately covers the same ground as a model card, but at the **system** level — i.e., the composite of Lexy’s prompts, model choices, and post-processing, not the underlying LLM weights (which Lexy does not train).

1. Intended purpose

Lexy assists employers (“deployers” in EU AI Act terms) in the recruiting process. The AI components:

- **Resume parsing** — extract structured fields from uploaded resumes.
- **Job-fit summarisation** — write a recruiter-facing summary comparing a candidate’s resume to the job description.
- **Interview generation** — generate role-specific interview questions from the JD + candidate background.
- **Interview conduct** — conduct a structured screening interview with the candidate, recording the conversation.
- **Interview scoring** — produce a structured assessment of the recorded interview against pre-defined competencies.

Lexy is **not** the decision-maker: every hire/reject decision is made by a named human recruiter or hiring manager in the customer’s organisation. Lexy’s outputs are inputs to a human-in-the-loop workflow.

This classifies Lexy as a high-risk AI system under EU AI Act Annex III §4 (“AI systems intended to be used for the recruitment or selection of natural persons”).

2. Provider and deployer

- **Provider** (EU AI Act Article 3(3)): Lexy Inc. (Delaware, USA).
- **Deployer** (Article 3(4)): each Customer using Lexy to fill its own open roles. Deployers have separate obligations under Article 26.

3. Model inventory

Function	Provider	Model	Version pinning	Training-data opt-out
Resume parsing	OpenAI	gpt-4o-mini	Pinned by API parameter	Zero-retention API contract
JD/fit summary	Anthropic	claude-3-5-sonnet	Pinned by API parameter	Zero-retention contract
Interview generation	OpenAI	gpt-4o	Pinned by API parameter	Zero-retention API contract
Interview conduct (voice + reasoning)	OpenAI	gpt-4o-realtime	Pinned by API parameter	Zero-retention API contract

Function	Provider	Model	Version pinning	Training-data opt-out
Interview scoring	Anthropic	claude-3-5-sonnet	Pinned by API parameter	Zero-retention contract

Lexy does **not** fine-tune or train on Customer Data. All inference calls are made with the provider’s zero-retention setting where available.

4. Inputs and outputs

Stage	Inputs	Outputs
Resume parsing	Resume text/PDF	Structured fields (name, contact, skills, work history)
Fit summary	Resume fields, JD text	Free-text summary + 0-100 fit score
Interview generation	JD, interview-plan template, resume highlights	Ordered list of 3-12 questions
Interview conduct	Live candidate audio + previous turns	Spoken follow-ups, transcript
Scoring	Transcript, JD, scoring rubric	Per-competency score + free-text rationale

The score and rationale outputs are surfaced to the recruiter alongside a prominent “AI-generated — review carefully” affordance.

5. Evaluated characteristics

Lexy evaluates candidates **only** on:

- Role-relevant skills, knowledge, and experience as expressed in the resume and during the interview.
- Communication clarity, structured thinking, and demonstrated reasoning in the candidate’s own answers.

Lexy does **not** evaluate, infer, or use:

- Race, ethnicity, national origin (beyond work authorisation, which the candidate self-reports).
- Age, religion, sexual orientation.
- Disability or family / marital status.
- Gender identity.
- Pitch, accent, or speech characteristics as a proxy for ability.
- Facial expressions, micro-expressions, or any video-based personality inference.

The interview system prompt explicitly forbids the AI from asking questions about any of these topics (see `artifacts/api-server/src/routes/interviews.ts`).

6. Known limitations and risks

- **LLM hallucination.** Outputs can be plausible but incorrect; the human-in-the-loop is the mitigation.

- **Score sensitivity to phrasing.** Two paraphrasings of the same answer may receive different scores. Mitigation: scoring is paired with free-text rationale so reviewers can sanity-check.
- **Language bias.** The underlying models are stronger in English than in other languages. Lexy supports multiple languages but accuracy in English is highest. Customers operating outside English should review outputs more carefully.
- **Disparate-impact risk.** AI hiring tools have historically shown bias against protected groups. Lexy mitigates by (a) decoupling demographics from the recruiter view, (b) forbidding the AI from evaluating demographic traits, and © providing an aggregate diversity dashboard with k-anonymity ≥ 5 so customers can detect imbalance. Customers using Lexy in NYC must additionally run an independent bias audit per Local Law 144; Lexy captures the underlying decision-log data today (`ai_decision_log` table) and provides an auditor-formatted CSV on request — a self-serve `/analytics/aedt-export` endpoint is a near-term follow-up.

7. Human oversight design

- Every AI score is displayed alongside a “Why this score?” rationale.
- No AI-only auto-reject: rejection requires a named recruiter to click “Reject” with a reason.
- Recruiters can override or hide any AI output without affecting the underlying candidate record.
- The audit log records every AI recommendation and every human decision so the chain of responsibility is reconstructible.

8. Post-market monitoring plan

- Weekly: review the rate of AI-recommended rejections by demographic bucket (where the customer has the aggregated demographics dashboard enabled). Flag any bucket below 0.8× the highest bucket as candidate for investigation (4/5ths rule).
- Monthly: sample 20 transcripts × score pairs and have a human reviewer rate inter-rater agreement with the AI.
- Quarterly: re-run the bias audit on a fresh sample for any customer hiring into NYC.
- On regulator request: the `ai_decision_log` table provides auditor-reproducible records per Local Law 144 (CSV export of the joined decision-log + demographics view with k-anonymity ≥ 5 is a near-term follow-up — the underlying data is captured today).

9. Change log

Version	Date	Change
1.0	2026-05-16	Initial draft